

LAB2: Local Methods and Cross Validation

Goal: to offer some familiarity with the k Nearest Neighbors (kNN) algorithm and get a practical understanding of the bias-variance trade-off using synthetic data.

This lab is divided into two parts depending of their level of complexity (**Beginner, Advanced**). Your goal is to complete entirely, at least, one of the two parts. Please note that a different notation can be used in the code as we used in the lectures.

PART I: Beginner

Warm up: Data Generation

1. Using the class `MixGauss.py` with appropriate parameters, produce a dataset with four classes and 30 samples per class: the classes must live in the 2D space and be centered on the corners of the unit square (0,0), (0,1), (1,1), (1,0), all with variance 0.3.
To run the example: `Xtr, Ytr = MixGauss([...])`
To visualise the data: `plt.scatter(Xtr[:,0], Xtr[:,1], 25, Ytr)`
`plt.show()`
2. Manipulate the data so to obtain a 2-class problem where data on opposite corners share the same class with labels +1 and -1. Hint: if you produced the data following the centers order provided above, you can use a mapping like
`Ytr = (Ytr % 2) * 2 - 1;`
3. Similarly generate a "test set" `Xte, Yte` drawn from the same distribution (start with 200 samples per class) and plot it using `plt.scatter()`.

The Geek part

The *k-Nearest Neighbors* algorithm (kNN) assigns to a test point the most frequent label of its k closest examples in the training set. Study the code of class `kNNClassify.py`.

4. Use `kNNClassify.py` on the 2-class data generated at step 1. Pick a "reasonable" number of neighbours k .
5. Evaluate the classification performance by comparing the estimated to the true outputs:
`sum(Ypred~=Yte)/Nt` %Nt number of test data
6. To visualise the separating function (and thus visualise what areas of the 2D plane are

associated with each class) you can use the provided `class separatingFkNN.py` and again plot the test points with `plt.scatter()`.

Analysis

So far we considered an arbitrary choice for k .

7. Perform a hold-out cross validation procedure on the available training data for a large range of candidate values for k (e.g. $k=1, 3, 5, \dots, 41$). Repeat the hold-out experiment for $\text{rep}=10$ time using at each iteration $p=30\%$ of the training set for validation. You can use the provided class `holdoutCVkNN.py` (check the code for an example use). Look at the training and validation errors for the different values of k . How would you now answer the question "what is the best value for k "? Note: for the parameters $\text{rep}=10$ and $p=0.3$, the hold-out procedure may be quite unstable.
8. How is the value affected by p (percentage of points held out) and number rep (number of repetitions e.g., 1, 5, 30, 50, 100) of the experiment? What does a large number of repetitions provide?
9. Apply the model obtained by cross validation (i.e., best k) to the test set (X_{te}) and see if there is an improvement on the classification error over the result of point 5 above.

PART II: Advanced

Advanced Analysis

10. *Dependence on training size:* Evaluate the performance as the size of the training set grows (varies), e.g., $n = \{3, 5, 20, 50, 100, 300, 500, \dots\}$. How would you choose a good range for k as n changes? Repeat the validation and test multiple times. What can you say about the stability of the solution/performances?
11. Try classifying more difficult datasets, generated through `MixGauss.py`, for instance, by increasing variance.
12. Modify the class `kNNClassify.py` to handle multi-class problems.
13. Modify the class `kNNClassify.py` to handle regression problems.