



# Machine Learning in Practice

## Crash Course on Machine Learning

August 22-24, 2016

**Arnaud Gotlieb and Valeriya Naumova**

Simula Research Laboratory AS

Simula School of Research and Innovation (SSRI)

# GOAL:

Advice on how to apply learning algorithms to different applications

## Some key aspects of this lecture:

- ▶ **No math!!!** But it could be much harder material to understand and use;
- ▶ Some aspects are debatable;
- ▶ Advice might not be applicable for novel machine learning research;
- ▶ Briefly.... to give you some time to play with the labs.

## Slides based on:

- ML Lecture by A. Ng, Stanford University
- Lectures and papers by P. Domingos, UC Washington
- Presentations by Scott Fortmann-Roe

# ML in Practice

- ▶ Understanding domain, prior knowledge, and goals
- ▶ Data integration, selection, cleaning, pre-processing, etc.
- ▶ Learning models
- ▶ Interpreting results
- ▶ Consolidating and deploying discovered knowledge
- ▶ Loop

**LEARNING**  
**=**  
**REPRESENTATION + EVALUATION + OPTIMISATION**

# Key aspects to remember

## ▶ It's generalisation that counts...


- ▶ Set some data aside from the beginning to test your estimator at the end
- ▶ Use cross-validation

## ▶ Data alone is not enough

- ▶ Every learner must embody some knowledge or assumption beyond the data it is given in order to generalise beyond it.
- ▶ **"No Free Lunch Theorem"**
- ▶ one of the key criteria for choosing a representation is which kinds of knowledge are easily expressed in it.
- ▶ **Remember: Machine Learning is not Magic!!!**

## ▶ Overfitting has many faces

- ▶ Decompose the generalisation error into bias and variance
- ▶ Use cross-validation, regularisation



**High Variance = Overfitting:**

- ▶ the model has too many parameters.

**High Bias = Underfitting:**

- ▶ the model is too rigid.

# Key aspects to remember

## ▶ Intuition fails in high dimensions

- ▶ Our intuitions, which come from a three- dimensional world, often do not apply in high-dimensional ones.
- ▶ Luckily most of the real-life data has a lower-dimensional representation

## ▶ Theoretical guarantees are not what they seem

## ▶ Feature engineering is the key

- ▶ Data pre-processing and feature extraction might be the most tedious work

## ▶ More data beats a cleverer algorithm

- ▶ The issue of scalability (time, memory and training set)

# Key aspects to remember

- ▶ **Learn many models not just one**

- ▶ Model ensembles: bagging, boosting..

- ▶ **Simplicity does not imply accuracy**

- ▶ **Representable does not imply learnable**

- ▶ Can it be represented?  Can it be learned?

- ▶ **Correlation does not imply causation**

- ▶ Diapers - Beer Example



# Getting Started on a Problem: Two Approaches

## Approach #1: Careful design.

- ▶ Spend a long term designing exactly the right features, collecting the right dataset, and designing the right algorithmic architecture.
- ▶ Implement it and hope it works.

**Benefit:** Nicer, perhaps more scalable algorithms. May come up with new, elegant, learning algorithms; contribute to basic research in machine learning.

## Approach #2: Build-and-fix.

- ▶ Implement something quick-and-dirty.
- ▶ Run error analyses and diagnostics to see what's wrong with it, and fix its errors.

**Benefit:** Will often get your application problem working more quickly. Faster time to market.



# Debugging Learning Algorithms

## Motivating Example

- ▶ Anti-spam. You carefully choose a small set of 100 words to use as features. (Instead of using all 50000+ words in English.)
- ▶ Bayesian logistic regression, implemented with gradient descent, gets 20% test error, which is unacceptably high.

$$\max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) - \lambda \|\theta\|^2$$

- ▶ What to do next?

# Fixing the Learning Algorithm

**Bayesian logistic regression:**

$$\max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) - \lambda \|\theta\|^2$$

Try improving algorithms in different ways:

- ▶ Getting more training examples.
  - ▶ Reduce the set of features.
  - ▶ Enlarge the set of features.
  - ▶ Use different features.
  - ▶ Run the optimiser (gradient descent) for some more iterations.
  - ▶ Choose a different optimisation algorithm.
  - ▶ Use a different regularisation term or constant value.
  - ▶ Try another learning algorithm (SVM).
- . . . some may be fixing problems you don't have.

**This approach might work, but it's very time-consuming, and largely a matter of luck whether you end up fixing what the problem really is.**

# Principled Analysis: Diagnostics

**First figure out what's going on.**

- ▶ Overfitting vs. Underfitting?
- ▶ Search error vs. Modelling error?
- ▶ Complex system: Find the most problematic component.

**Trivial but vital:**

- ▶ Visualise the data. (Plot or view frequent patterns.)
- ▶ Start with simple things.

# Diagnostic for Bias vs Variance

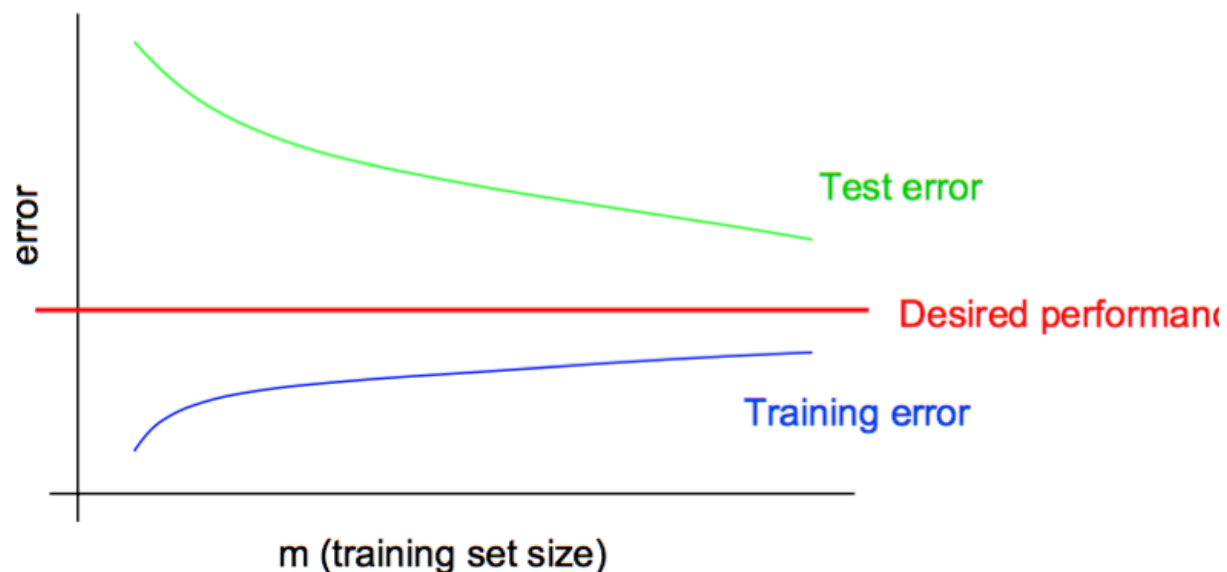
Suppose you suspect the problem is either:

- Overfitting (high variance).
- Too few features to classify spam (high bias).


Diagnostic:

- **Variance**: Training error will be much lower than test error.
- **Bias**: Training error will also be high.

Typical learning curve for  
high variance



Typical learning curve for  
high bias



# Diagnostics Tell You What to Try Next

**Bayesian logistic regression:**

$$\max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) - \lambda \|\theta\|^2$$

Try improving algorithms in different ways:

- ▶ Getting more training examples. **Fixes high variance**
  - ▶ Reduce the set of features. **Fixes high variance**
  - ▶ Enlarge the set of features. **Fixes high bias**
  - ▶ Use different features (email header). **Fixes high bias**
  - ▶ Run the optimiser (gradient descent) for some more iterations.
  - ▶ Choose a different optimisation algorithm.
  - ▶ Use a different regularisation term or constant value.
  - ▶ Try another learning algorithm (SVM).
- . . . some may be fixing problems you don't have.

# Optimisation Algorithm Diagnostics

- ▶ Bias vs. variance is one common diagnostic.
- ▶ For other problems, it's usually up to your own ingenuity to construct your own diagnostics to figure out what's wrong.
- ▶ **Another example:**
  - Bayesian logistic regression gets 2% error on spam, and 2% error on non-spam. (Unacceptably high error on non-spam.)
  - SVM using a linear kernel gets 10% error on spam, and 0.01% error on non-spam. (Acceptable performance.)
  - But you want to use logistic regression, because of computational efficiency, etc.
- ▶ What to do next?

# Search vs Modelling Error

## Search Error:

- ▶ the optimiser fails to find the best parameters
- ▶ . . . a problem with the optimiser.

## Modelling Error:

- ▶ the best parameters do not lead to the best performance.
- ▶ . . . a problem with the objective function.

## Consider:

- ▶ Will more iterations help? Is the algorithm (gradient descent for logistic regression) converging?
- ▶ When can two learners help to diagnose the problem?
- ▶ Are you optimising the right function?
- ▶ Correct value of the regularisation parameter?



# Diagnostics tell you what to try next

## Bayesian logistic regression:

$$\max_{\theta} \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}, \theta) - \lambda \|\theta\|^2$$

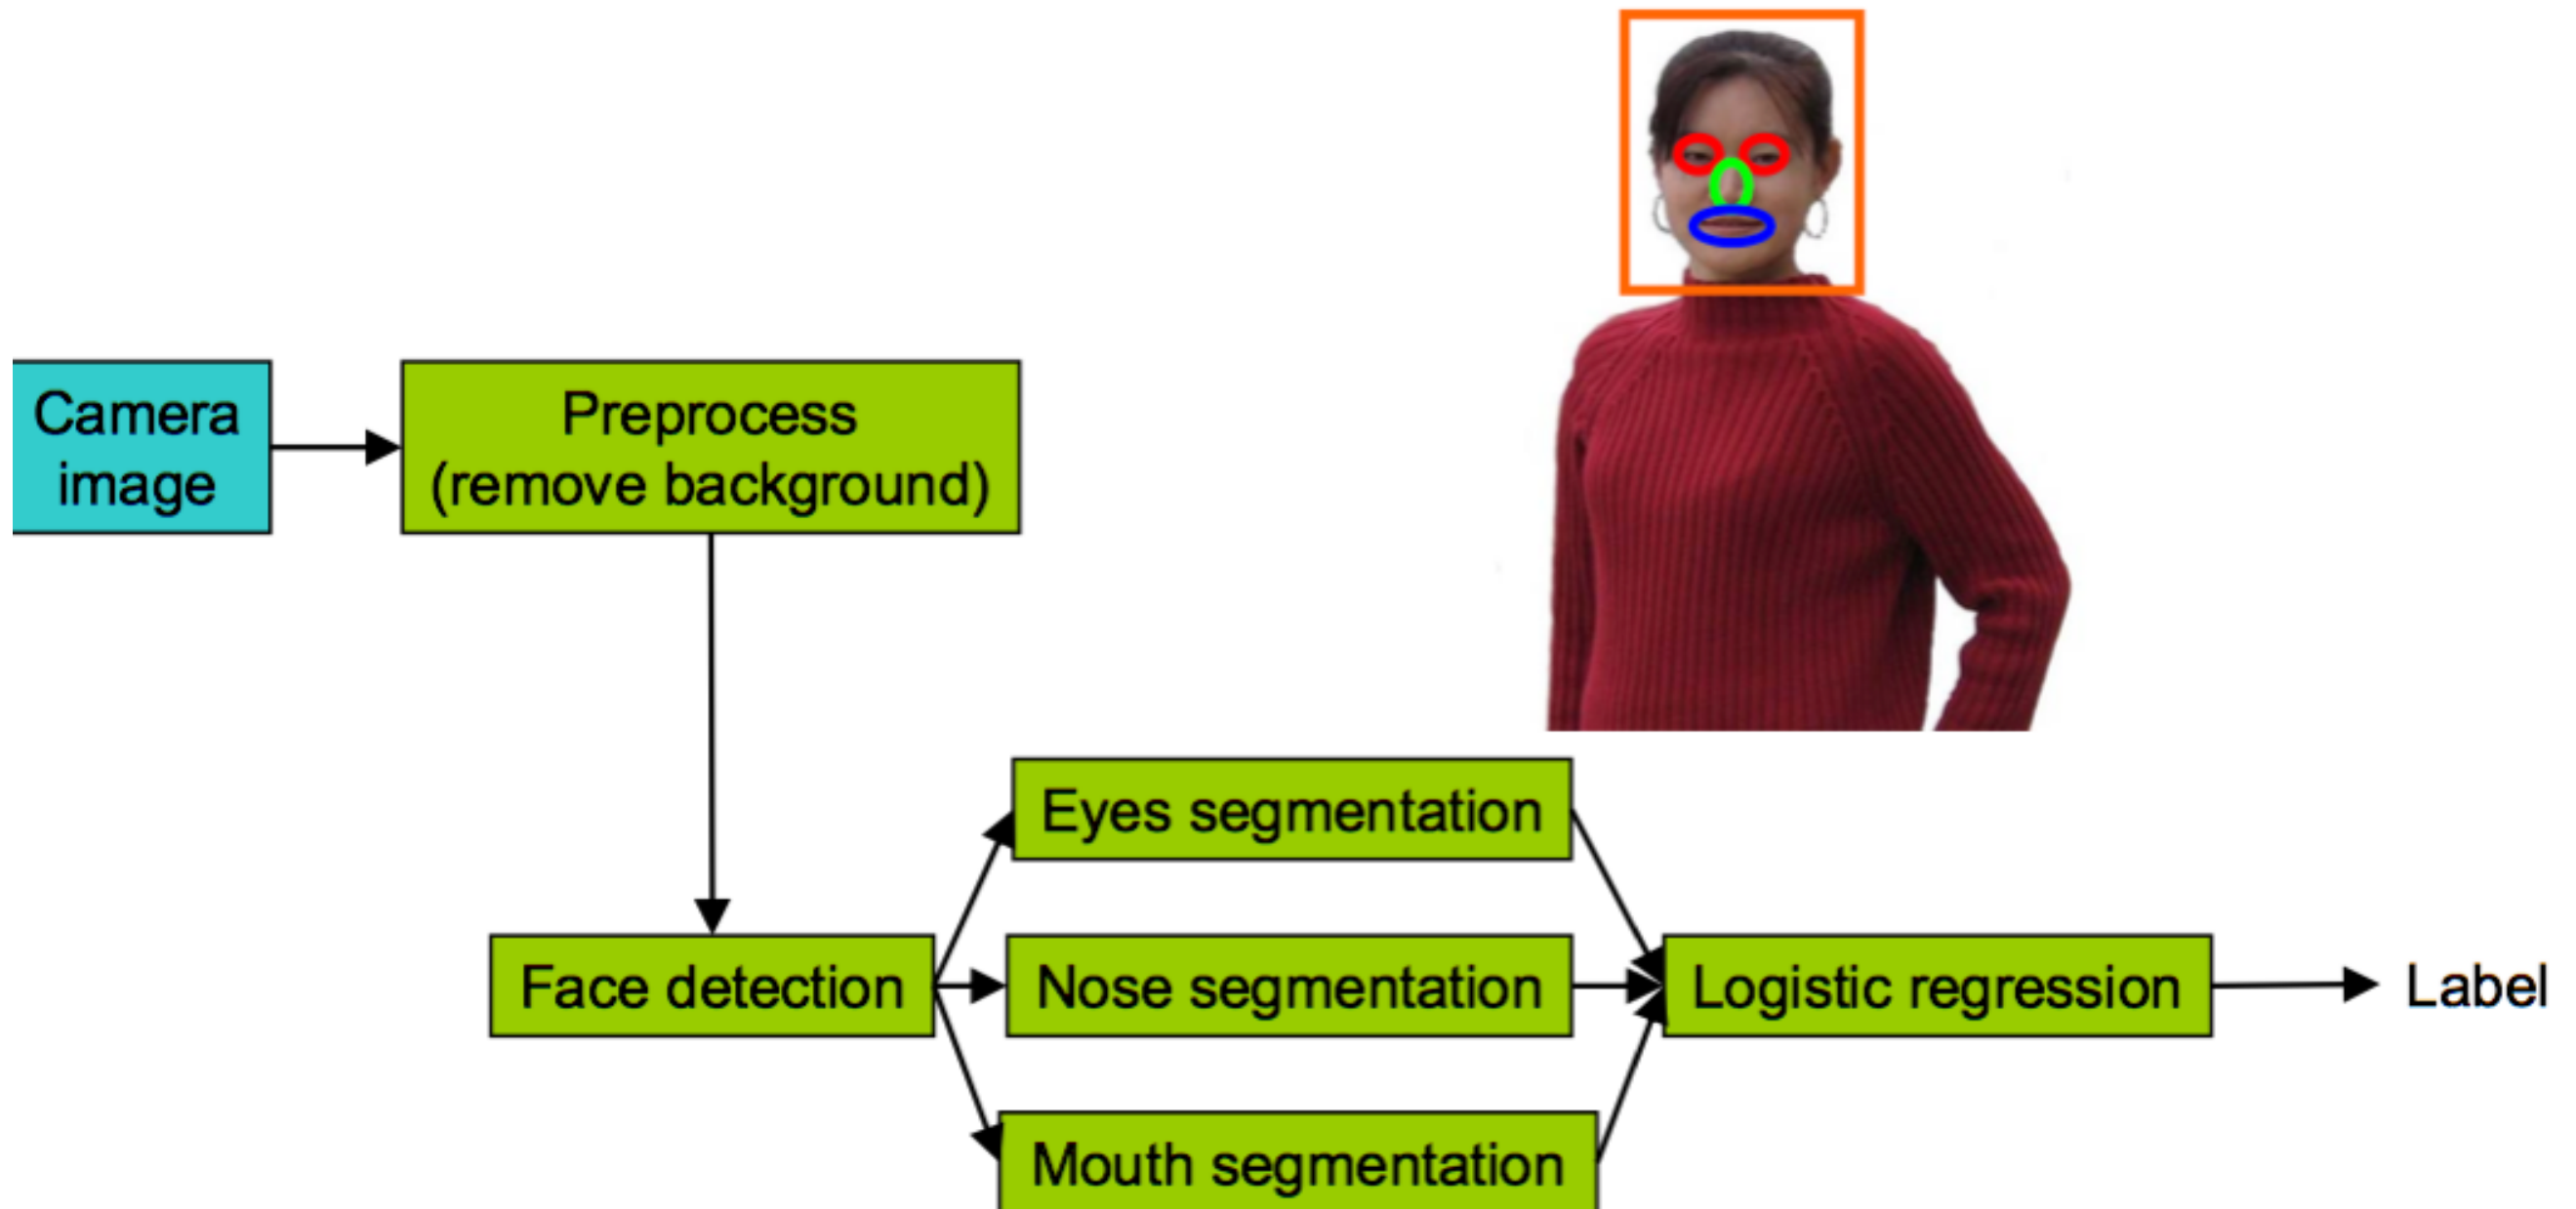
Try improving algorithms in different ways:

- ▶ Getting more training examples. **Fixes high variance**
  - ▶ Reduce the set of features. **Fixes high variance**
  - ▶ Enlarge the set of features. **Fixes high bias**
  - ▶ Use different features (email header). **Fixes high bias**
  - ▶ Run the optimiser (gradient descent) for some more iterations. **Fixes opt algorithm**
  - ▶ Choose a different optimisation algorithm. **Fixes opt algorithm**
  - ▶ Use a different regularisation term or constant value. **Fixes opt objective**
  - ▶ Try another learning algorithm (SVM). **Fixes opt objective**
- . . . some may be fixing problems you don't have.

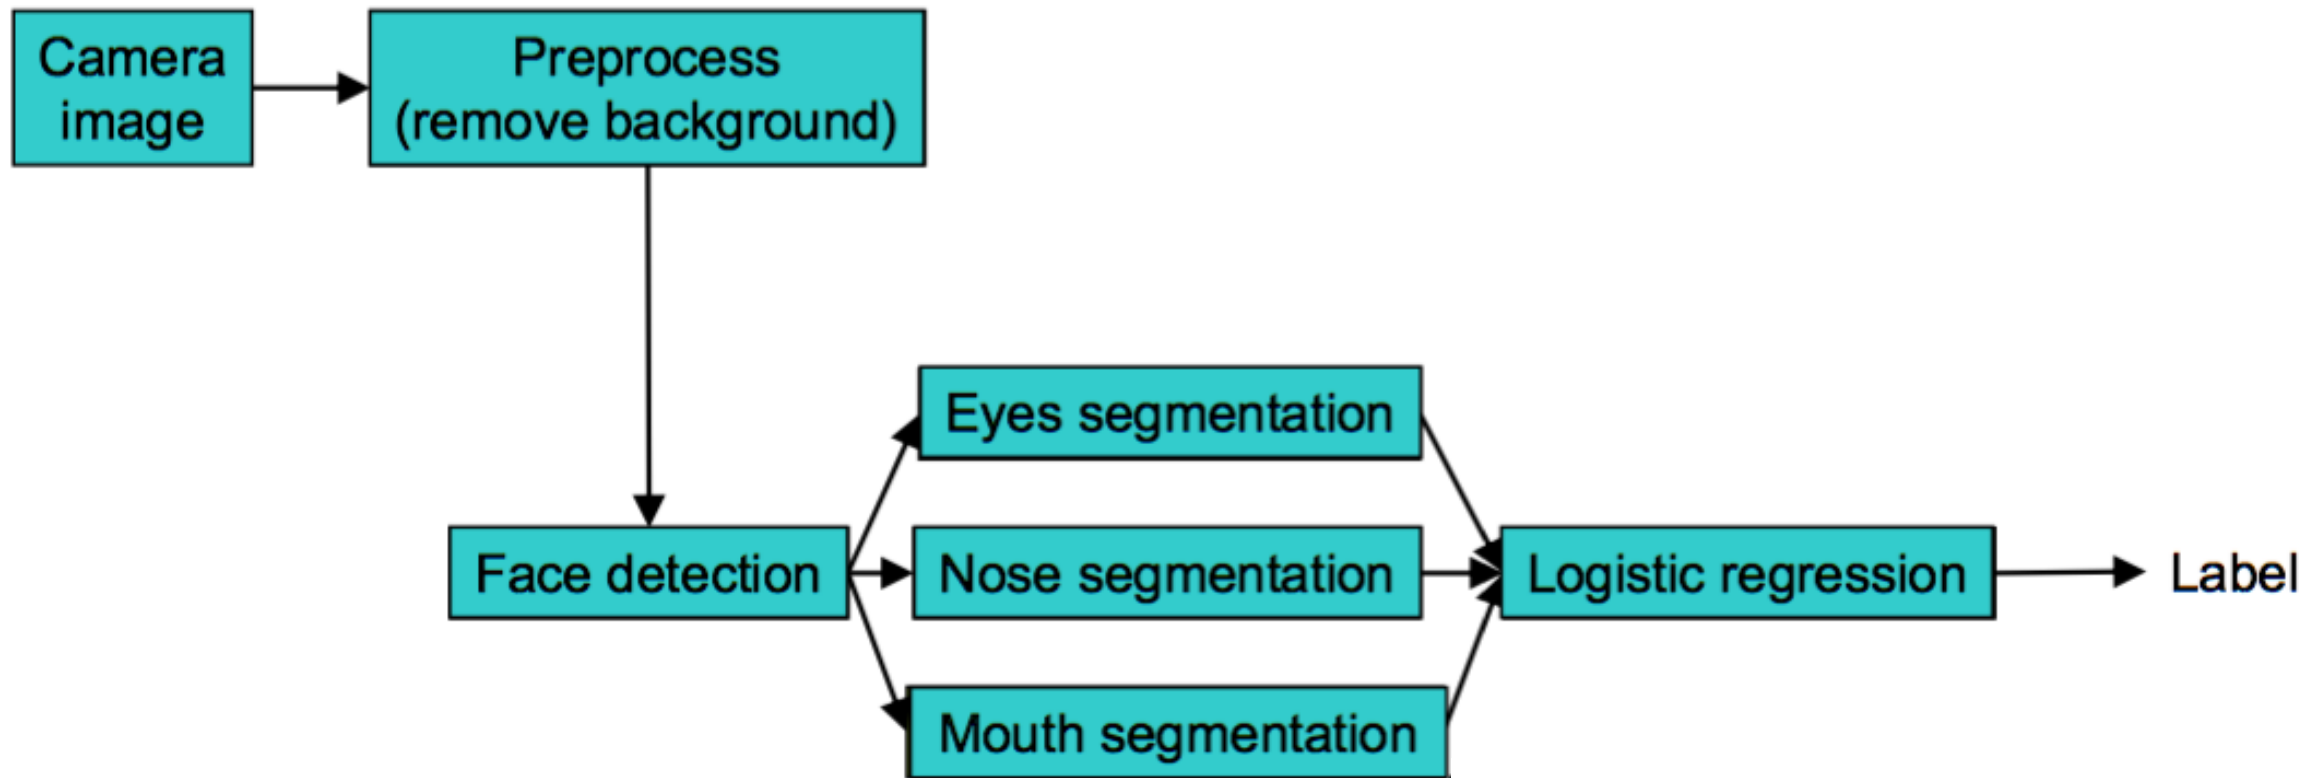
# Error Analysis

Machine learning “pipeline” consists of many learning algorithms.

Example: face recognition from images.



# Error Analysis



How much error is attributable to each of the components?

Plug in ground-truth for each component, and see how accuracy changes.

Conclusion: Most room for improvement in face detection and eyes segmentation.

Component	Accuracy
Overall System	85 %
Preprocess (remove background)	85 %
Face detection	91 %
Eyes segmentation	95 %
Nose segmentation	96 %
Mouth segmentation	97 %
Logistic Regression	100 %

# Complex Systems

## Error Analysis:

- ▶ Compares the best possible vs. current accuracy.
- ▶ Provide more and more golden truth data as part of the input.
- ▶ Find the component where the jump in accuracy is the highest.

## Ablative Analysis:

- ▶ Compares some baseline vs. current accuracy.
- ▶ Switch off more and more components.
- ▶ Find the component where the loss in accuracy is the highest.